

Multi-agent systems - implementation of reinforcement learning and Markov decision processes using FrozenLake from gym library


Guy Cohen

Ester Vaknin

Almog Klein

Afeka College of Engineering

April 07, 2022



In this report we will present all the steps we implemented in the code and describe the results we got and draw conclusions.

שלב 1: עליכם לבחור משימה עבור סוכן בודד, שמקיימת תנאי MDP. נא הציגו את המשימה כתהליך קבלת החלטות מרקובי סופי. יש להגדיר באופן מפורש את המצבים (states), פעולות האפשריות בכל מצב, הסתברויות מעבר (transition probabilities) והתשלומים (expected rewards). כמו כן, נא ציין במפורש את הכללים המיוחדים שיש בסביבה שהגדרתם.

Frozen lake involves crossing a frozen lake from Start(S) to Goal(G) without falling into any Holes(H) by walking over the Frozen(F) lake. The agent may not always move in the intended direction due to the slippery nature of the frozen lake. The agent takes a 1-element vector for actions. The action space is (dir), where dir decides direction to move in which can be:

| Action Space | |
|--------------|--------|
| NO. | Action |
| 0 | West |
| 1 | south |
| 2 | East |
| 3 | North |



is_slippery: If True, will move in intended direction with probability of 1/3 else will move in either perpendicular direction with equal probability of 1/3 in both directions.

The observation is a value representing the agent’s current position as `current_row * nrows + current_col` (where both the row and col start at 0).

Reward schedule:

- Reach goal(G): +1 (terminal state)
- Reach hole(H): 0 (terminal state)
- Reach frozen(F): 0

$$G_t \doteq R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \cdots = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1}$$

$$\Pr\{R_{t+1} = r, S_{t+1} = s' \mid S_0, A_0, R_1, \dots, S_{t-1}, A_{t-1}, R_t, S_t, A_t\},$$

$$p(s', r | s, a) = \Pr\{S_{t+1}=s', R_{t+1}=r \mid S_t=s, A_t=a\}.$$

$$p(s'|s, a) = \Pr\{S_{t+1}=s' \mid S_t=s, A_t=a\} = \sum_{r \in \mathcal{R}} p(s', r|s, a),$$

$$r(s, a, s') = \mathbb{E}[R_{t+1} \mid S_t = s, A_t = a, S_{t+1} = s'] = \frac{\sum_{r \in \mathcal{R}} rp(s', r | s, a)}{p(s' | s, a)}.$$

Transition matrix =

Reward vector =

שלב 2: עבור מדיניות π (policy) מסוימת (לבחירתכם) נא חישוב ערכי $V^\pi(s)$ לכל מצב s באמצעות משוואות בלמן. נא רשמו את המערכת ואת הפתרון. כמו כן, עבור אותה מדיניות π נא חשבו $V^\pi(s)$ באמצעות Value Iteration Algorithm. הוסיפו קטע קוד רלוונטי, ציינו מהם תנאי העצירה ורשמו כמות האיטרציות שהתבצעו. הסיקו מסקנות.

the *value* of a state s under a policy π , denoted $v_\pi(s)$, is the expected return when starting in s and following π thereafter. For MDPs, we can define $v_\pi(s)$ formally as

$$v_\pi(s) = \mathbb{E}_\pi[G_t \mid S_t = s] = \mathbb{E}_\pi \left[\sum_{k=0}^{\infty} \gamma^k R_{t+k+1} \mid S_t = s \right], \quad (3.10)$$

$$= \sum_a \pi(a|s) \sum_{s',r} p(s',r|s,a) \left[r + \gamma v_\pi(s') \right]$$

```

Initialize array  $V$  arbitrarily (e.g.,  $V(s) = 0$  for all  $s \in S^+$ )
Repeat
   $\Delta \leftarrow 0$ 
  For each  $s \in S$ :
     $v \leftarrow V(s)$ 
     $V(s) \leftarrow \max_a \sum_{s',r} p(s',r|s,a) [r + \gamma V(s')]$ 
     $\Delta \leftarrow \max(\Delta, |v - V(s)|)$ 
until  $\Delta < \theta$  (a small positive number)

Output a deterministic policy,  $\pi$ , such that
 $\pi(s) = \operatorname{argmax}_a \sum_{s',r} p(s',r|s,a) [r + \gamma V(s')]$ 

```

Calculation using Value Iteration Algorithm:

The number of iterations performed, policy and value function:

```

value iteration for 60 iterations
agent succeeded to reach goal 730117 out of 1000000 Episodes

policy:

[['West(<)' 'North(^)' 'West(<)' 'North(^)']
 ['West(<)' 'West(<)' 'West(<)' 'West(<)']
 ['North(^)' 'south(v)' 'West(<)' 'West(<)']
 ['West(<)' 'East(>)' 'south(v)' 'West(<)']]

value func vector:

[[0.06888624 0.06141117 0.07440763 0.05580502]
 [0.09185097 0.          0.11220727 0.          ]
 [0.14543392 0.24749561 0.29961676 0.          ]
 [0.          0.37993504 0.63901974 0.          ]]

```

stopping criteria and bellman equations:

```

Stop conditions:
delta < theta : delta = 9.568972661466724e-07

value matrix:

[[0.06888703374487129, 0.06664451425610295, 0.06664451425610295, 0.059755096254593354]
 [0.03908922434013946, 0.04298816225802422, 0.04074564276925588, 0.061411514683709784]
 [0.0744078231426156, 0.06882703782948876, 0.07272597574737354, 0.05748714779969938]
 [0.03906379537401382, 0.03906379537401382, 0.033483010060886984, 0.05580530040445731]
 [0.09185133639510051, 0.07118546448064662, 0.06429604647913703, 0.048221161830417386]
 [0.0, 0.0, 0.0, 0.0]
 [0.1122073190157469, 0.08988502867217658, 0.1122073190157469, 0.022322290343570327]
 [0.0, 0.0, 0.0, 0.0]
 [0.07118546448064662, 0.11787885805557752, 0.10180397340685787, 0.145434147971541]
 [0.15761068697605268, 0.24749571564822925, 0.20386554108354613, 0.13351520323685973]
 [0.29961678710056533, 0.2659546067272056, 0.22536810360967094, 0.10791086386425411]
 [0.0, 0.0, 0.0, 0.0]
 [0.0, 0.0, 0.0, 0.0]
 [0.18822919590226395, 0.3056864356476808, 0.37993511913857514, 0.2659546067272056]
 [0.39557146431985735, 0.6390197689810141, 0.6149242852418212, 0.5371988744168794]
 [0.0, 0.0, 0.0, 0.0]]

```

At this stage, we presented the Bellman Equation and calculated $V^\pi(s)$.

The policy we chose is: $\pi(a|s) = \frac{1}{4}$

We found that the number of iterations we received is 60 iterations and the stopping condition received is $9.56 \times 10^{-7} < 1 \times 10^{-6}$. We assume that the algorithm converged to these results due that this game consists of a number of basic and simple operations.

שלב 3: השתמשו ב- Policy Iteration Algorithm למציאת מדיניות אופטימאלית. הציגו את המדיניות האופטימאלית. מהם ערכי $V(s)$ עבור מדיניות אופטימאלית? הוסיפו קטע קוד רלוונטי, ציינו מהם תנאי העצירה ורשמו כמות האיטרציות שהתבצעו. הסיקו מסקנות.

This is an *optimal policy*. Although there may be more than one, we denote all the optimal policies by π_* . They share the same state-value function, called the *optimal state-value function*, denoted v_* , and defined as

$$v_*(s) = \max_{\pi} v_{\pi}(s), \quad (3.13)$$

for all $s \in \mathcal{S}$.

```
1. Initialization
    $V(s) \in \mathbb{R}$  and  $\pi(s) \in \mathcal{A}(s)$  arbitrarily for all  $s \in \mathcal{S}$ 

2. Policy Evaluation
   Repeat
      $\Delta \leftarrow 0$ 
     For each  $s \in \mathcal{S}$ :
        $v \leftarrow V(s)$ 
        $V(s) \leftarrow \sum_{s',r} p(s',r|s,\pi(s)) [r + \gamma V(s')]$ 
        $\Delta \leftarrow \max(\Delta, |v - V(s)|)$ 
   until  $\Delta < \theta$  (a small positive number)

3. Policy Improvement
   policy-stable  $\leftarrow$  true
   For each  $s \in \mathcal{S}$ :
      $a \leftarrow \pi(s)$ 
      $\pi(s) \leftarrow \arg \max_a \sum_{s',r} p(s',r|s,a) [r + \gamma V(s')]$ 
     If  $a \neq \pi(s)$ , then policy-stable  $\leftarrow$  false
   If policy-stable, then stop and return  $V$  and  $\pi$ ; else go to 2
```

$$\pi_0 \xrightarrow{E} v_{\pi_0} \xrightarrow{I} \pi_1 \xrightarrow{E} v_{\pi_1} \xrightarrow{I} \pi_2 \xrightarrow{E} \dots \xrightarrow{I} \pi_* \xrightarrow{E} v_*$$

```
entered stop criteria after 170 iterations with delta< theta : 8.88967639528504e-07 < 1e-06

*policy:

[['West(<)' 'North(^)' 'West(<)' 'North(^)']
 ['West(<)' 'West(<)' 'West(<)' 'West(<)']
 ['North(^)' 'South(v)' 'West(<)' 'West(<)']
 ['West(<)' 'East(>)' 'South(v)' 'West(<)']]
value func vector (*policy):

[[0.068888651 0.06141136 0.07440774 0.055805 ]
 [0.09185115 0.          0.11220731 0.          ]
 [0.14543402 0.24749566 0.29961679 0.          ]
 [0.          0.37993506 0.63901975 0.          ]]
```

Stopping criteria is when delta (the change of the value function between every iteration) is smaller than $\theta = 10^{-6}$. The number of iterations till the stop criteria occurred is 170. We can see that the optimal policy is the same as the policy we got in section 2. This can be explained by the simple game rules and the low number of states and actions. Also, the method we used in section 2 and section 3 converged to optimal policy (maybe local

to the corresponding algorithm for action values:

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha [R_{t+1} + \gamma Q(S_{t+1}, A_{t+1}) - Q(S_t, A_t)].$$

```
finished with 40114295 Steps for 1500000 episodes

***** Final Q-Table *****

[[0.03734557 0.0444739 0.03721523 0.03738171]
 [0.02620652 0.0293202 0.02664275 0.04329033]
 [0.06616942 0.04220908 0.04099227 0.04014203]
 [0.02793239 0.02612291 0.02664551 0.03556698]
 [0.06262842 0.04443403 0.04565074 0.03224161]
 [0. 0. 0. 0. ]
 [0.06894727 0.05949653 0.08551704 0.00821178]
 [0. 0. 0. 0. ]
 [0.06334845 0.06859887 0.07197902 0.13601024]
 [0.1336951 0.21425342 0.13819887 0.10641599]
 [0.25940121 0.18075051 0.14643228 0.12346345]
 [0. 0. 0. 0. ]
 [0. 0. 0. 0. ]
 [0.15905654 0.24654516 0.26742412 0.22370732]
 [0.38737202 0.55471987 0.47914216 0.4486898 ]
 [0. 0. 0. 0. ]]
```

Till this section we only used value function to measure the agent learning, and in this section, we will use Q learning. We chose to initialize the (16, 4) Q table with 0 as can be seen in the screenshot. We run the algorithm 1,500,000 episodes, which in them the agent took 40,114,295 steps in the game (26.5 in average for episode). We can see in the final Q table that there are 5 states that have not been learned. Those are the 4 holes in the grid (mentioned in section 1) and the goal state, meaning all the terminal states of the game will have Q table values as we initialize them. Another fact is as we chose bigger state id (meaning closer to the goal state) the values in the table rows are bigger, as expected for the agent learning.

שלב 6: בחלק זה נא השווא תוצאות של Q-learning algorithm עם/ללא eligibility traces. ניתן להציג את ההשוואה בצורה ויזואלית באמצעות גרפים של ההתכנסות. הסיקו מסקנות.


That is, for an episode starting at time step 0 and terminating at step T , for all $s \in \mathcal{S}$:

$$V_{t+1}(s) = V_t(s), \quad \forall t < T,$$


$$V_T(s) = V_{T-1}(s) + \sum_{t=0}^{T-1} \Delta_t(s),$$

$$\begin{aligned} G_t^\lambda &= \sum_{n=1}^{\infty} G_t^{(n)} (1 - \lambda_{t+n}) \prod_{i=t+1}^{t+n-1} \lambda_i \\ &= \sum_{k=t+1}^{T-1} G_t^{(k-t)} (1 - \lambda_k) \prod_{i=t+1}^{k-1} \lambda_i + G_t \prod_{i=t+1}^{T-1} \lambda_i. \end{aligned}$$

Learning progress for Q-Learning without eligibly traces



Learning progress for Q-Learning with eligibly traces




```

***** Q-Table *****
[[0.42611206 0.5146688 0.68378067 0.44952703]
 [0.56614517 0.73793606 0.59553851 0.54093271]
 [0.57281976 0.69919344 0.60728218 0.57057674]
 [0.81303196 0.6756013 0.62162979 0.58979903]
 [0.634719 0.78785183 0.55941805 0.54936696]
 [1. 1. 1. 1. ]
 [0.78182352 0.86066875 0.69403824 0.78128738]
 [1. 1. 1. 1. ]
 [0.64922002 0.60296072 0.85199473 0.62831113]
 [0.73994143 0.73854687 0.76214952 0.78406708]
 [0.95535285 0.76897928 0.78364447 0.77079921]
 [1. 1. 1. 1. ]
 [1. 1. 1. 1. ]
 [0.77830515 0.80920183 0.8069891 1.01877484]
 [0.8214491 1.39593189 0.98599238 1.05696981]
 [1. 1. 1. 1. ]]

***** ET-Table *****
[[3.96570448e-075 2.80194362e-006 9.07884771e-001 5.66107658e-093]
 [8.66966374e-230 1.12603384e-001 7.47901769e-135 3.34971326e-086]
 [8.34818567e-205 3.40885196e-005 3.00100081e-077 1.25156559e-260]
 [3.13163796e-008 4.91977510e-114 2.22656386e-195 9.88131292e-324]
 [2.00696597e-007 1.87816660e+000 8.96780076e-046 2.34654535e-068]
 [0.00000000e+000 0.00000000e+000 0.00000000e+000 0.00000000e+000]
 [1.18603504e-085 1.20216726e-012 9.88131292e-324 6.05155199e-086]
 [0.00000000e+000 0.00000000e+000 0.00000000e+000 0.00000000e+000]
 [1.38776219e-267 9.88131292e-324 1.36446600e+000 9.88131292e-324]
 [2.19662118e-292 1.01678023e-301 1.23393555e-206 1.00835464e-008]
 [1.42457656e-014 2.28826792e-207 9.88131292e-324 9.88131292e-324]
 [0.00000000e+000 0.00000000e+000 0.00000000e+000 0.00000000e+000]
 [0.00000000e+000 0.00000000e+000 0.00000000e+000 0.00000000e+000]
 [9.88131292e-324 9.88131292e-324 9.88131292e-324 7.93458946e-071]
 [9.88131292e-324 1.37865940e-034 1.37810377e-210 4.10887751e-210]
 [0.00000000e+000 0.00000000e+000 0.00000000e+000 0.00000000e+000]]

```

The method we used to measure the learning process is every (episodes//STEPS) episodes of the learning we took the existing policy and played 500 episodes in checked how many time the agent won.

The first graph is Q learning without eligibility traces (ET). The second is with ET. We can see that the learning without ET has fewer measures which the average reward is close to 0, than the with ET. The average value of both graphs is the same (maybe a little bit bigger for without ET). but the max values of the graphs are bigger for the with ET learning. The average win percentage is low (around 4%) for both learnings. probably because of the randomness in the environment as mentioned in section 3.